

MnnFast : A Fast and Scalable System Architecture for Memory-Augmented Neural Networks

Hanhwi Jang

46th International Symposium on Computer Architecture (ISCA), 2019

Presenter: Jaeyoon Lee

<http://esoc.hanyang.ac.kr/people/<userlink>>

March 26, 2020



Neural Network Acceleration Study Season #2

Contents of presentation

- **Introduction**

- Memory network(MemNNs)
- Bag-of-Words(BoW) models
- Computational steps of MemNN

- **Motivation**

- Three major performance problems of MemNN

- **MnnFast**

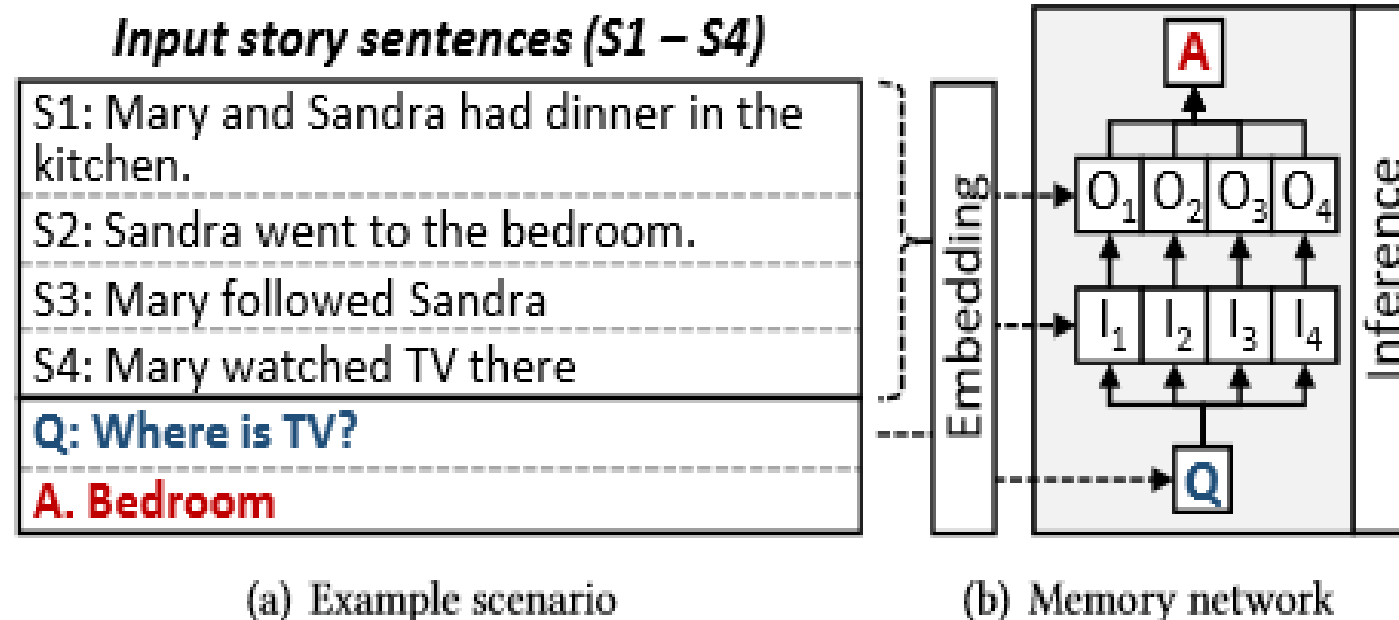
- Three ways to overcome performance problems of MemNN

- **Implementation & Evaluation**

- **Conclusion**

Introduction-Memory Network

- **Context-aware information processing model, MemNNs.**
 - Similar to the human's working memory
 - Values are stored into memory components



Introduction-BoW models

- Bag-of-Words model enable text data to be represented in the form which can be processed by ML algorithms

The Bag of Words Representation

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
would	1
whimsical	1
times	1
sweet	1
satirical	1
adventure	1
genre	1
fairy	1
humor	1
have	1
great	1
...	...

Introduction-Word Embedding

- Word Embedding converts words into meaningful vectors
- After **BoW model**, sentence become $(1 \times V)$ vector.

V : size of dictionary

- **Embedding operation** converts sentence vector $(1 \times V)$ into internal states $(1 \times ed)$

ed : embedding dimension

Introduction-Computational steps of MemNN

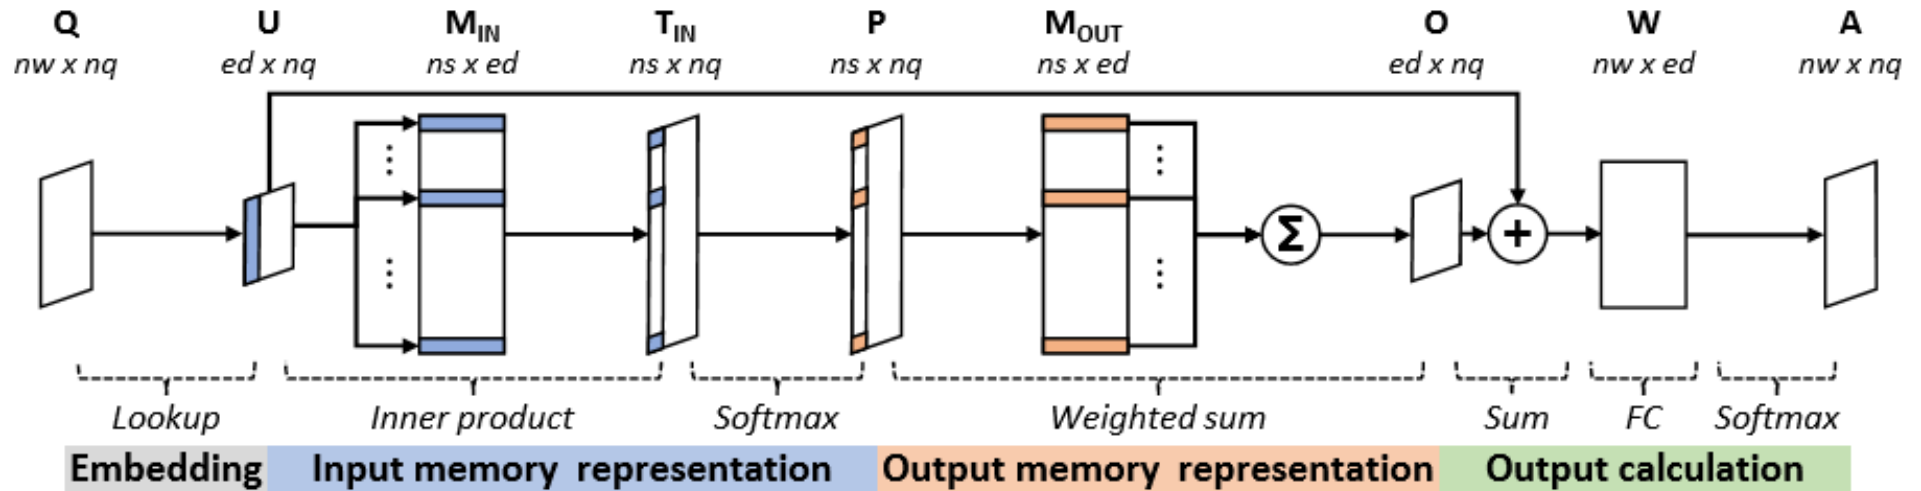
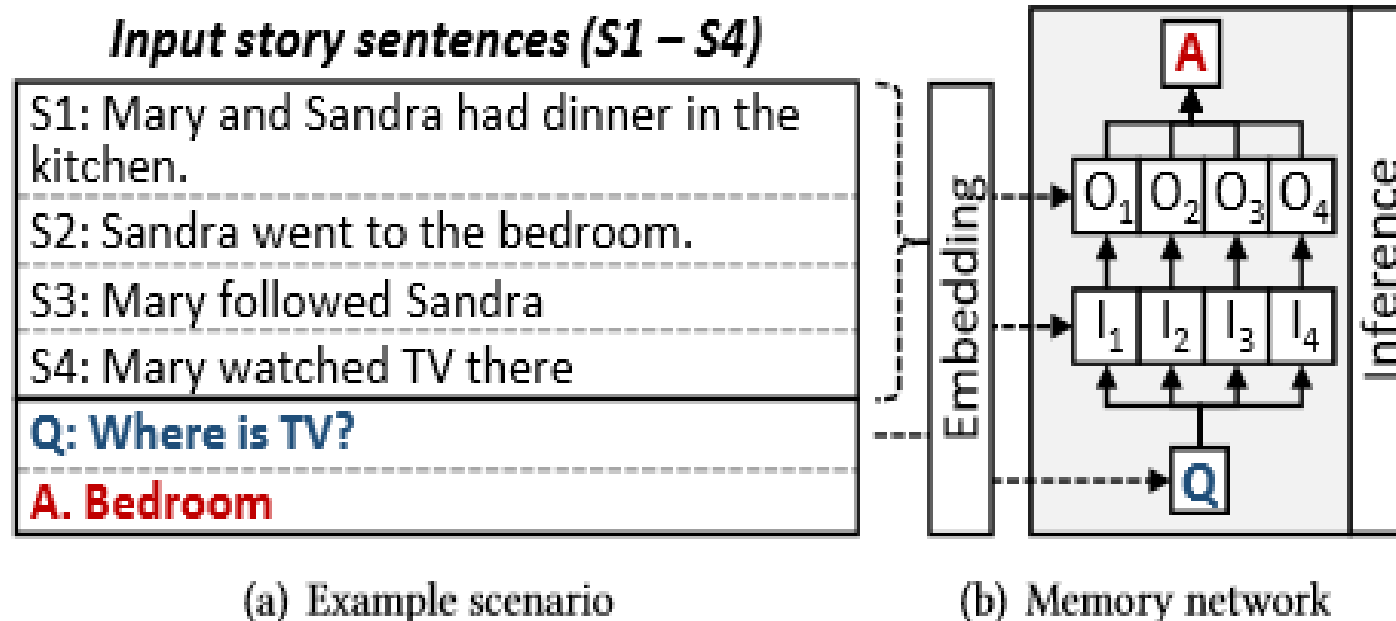


Figure 2: Computational steps of memory networks (MemNN). MemNN consists of embedding, input memory representation, output memory representation and output calculation. nw is the maximum number of words in a sentence. nq and ns are the number of questions and given story sentences, respectively. ed is the embedding dimension.

nw : max number of words in sentence
 nq : number of questions
 ed : embedding dimension
 ns : number of story sentence

Motivation

- **To improve MemNNs, we have to increase the size of memory.**
 - Requires a fast and scalable computer infrastructure.
 - Current system **does not provide enough scalability.**



Motivation- performance problem of MemNN

1. *High Memory BandWidth Consumption*

- Data do not fit into the cache
→ Data spills to DRAM memory
- Increased number of DRAM accesses degrades performance




Figure 3: Limited scalability due to memory bandwidth bottleneck. The speedup results of each channel configuration are normalized to the corresponding single-thread results.

Motivation- performance problem of MemNN


2. Heavy Computation

- Certain phases of MemNN consist of a large number of compute-intensive operation.
 - ex) Matrix inner product, weighted sum, softmax

Motivation- performance problem of MemNN

3. *Cache Contention*

- MemNN can suffer from **cache conflicts** because of shared cache
- *Inference operation* need to keep necessary data in shared cache
 - ↔ *Embedding operation* results in polluting shared cache



MnnFast-Column based Algorithm




Figure 2: Computational steps of memory networks (MemNN). MemNN consists of embedding, input memory representation, output memory representation and output calculation. nw is the maximum number of words in a sentence. nq and ns are the number of questions and given story sentences, respectively. ed is the embedding dimension.

MnnFast-Column based Algorithm

1. Baseline MemNN

Intermediate vector(T_{IN} , P_{exp} , P)

-proportionate to ns

-each vector's size is **800MB**

-spilled to DRAM

$$o = \sum_i \text{softmax}(u \times m_i^{IN}) m_i^{OUT}$$

$$= \sum_i \frac{e^{u \times m_i^{IN}} m_i^{OUT}}{\sum_j e^{u \times m_j^{IN}}}$$

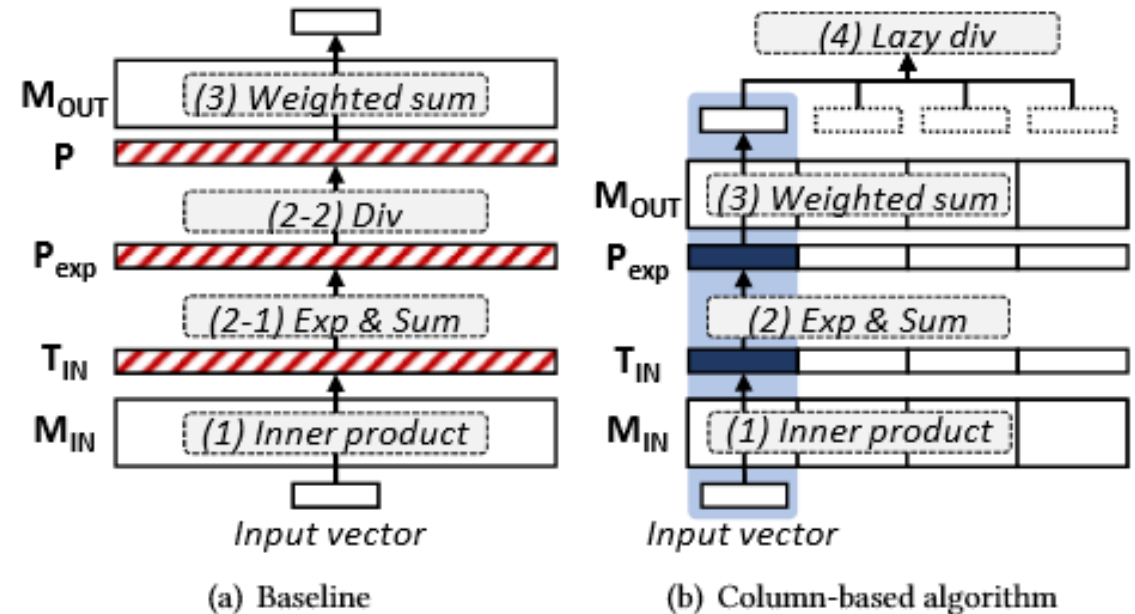


Figure 5: Dataflow comparison between the baseline and the column-based algorithm.

MnnFast-Column based Algorithm

2. Column-based Algorithm

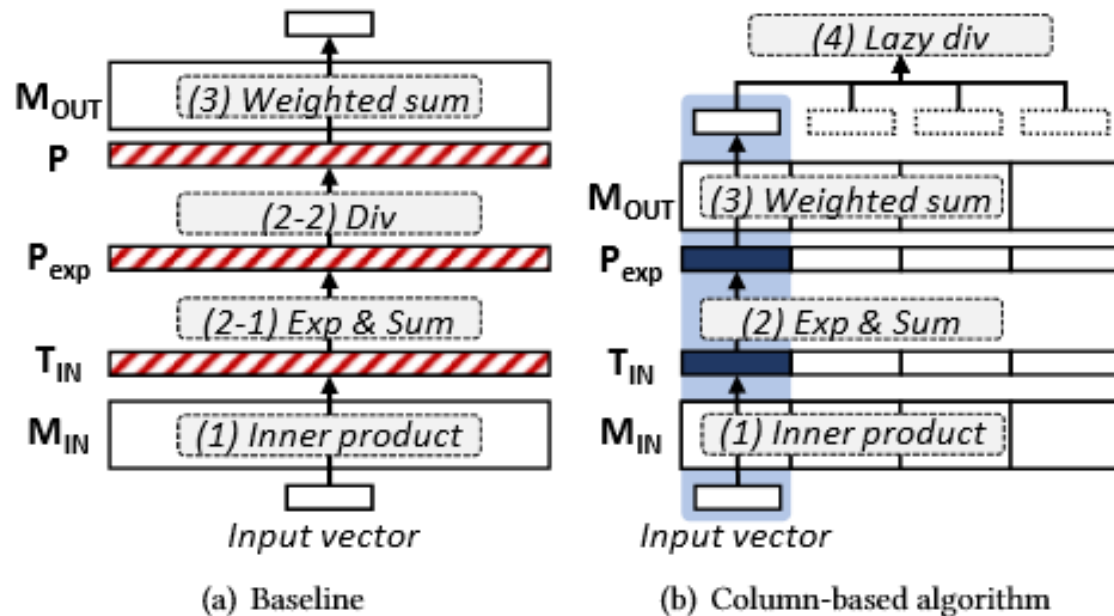


Figure 5: Dataflow comparison between the baseline and the column-based algorithm.


Lazy softmax

- compute Softmax's division at last
- same results as the baseline


$$o = \sum_i \text{softmax}(u \times m_i^{IN}) m_i^{OUT}$$

$$= \frac{1}{\sum_j e^{u \times m_j^{IN}}} \sum_i e^{u \times m_i^{IN}} m_i^{OUT}$$


MnnFast-Column based Algorithm




MnnFast-Column based Algorithm



MnnFast-Column based Algorithm



MnnFast-Column based Algorithm



MnnFast-Column based Algorithm

- By doing Column-based Algorithm,
 - 1) Reduce the size of temporary data to fit those into the on-chip cache
Column-based MemNN can load those memory into the cache. This leads to the capability of streaming.
 - 2) Reduce the amount of computation(softmax's division operation)
 - 3) Column-based MemNN can partition each layer into multiple sub-layers.

MnnFast-Zero Skipping

- In vector p , only small number of words and sentences are correlated with given question.




Figure 2: Computational steps of memory networks (MemNN). MemNN consists of embedding, input memory representation, output memory representation and output calculation. nw is the maximum number of words in a sentence. nq and ns are the number of questions and given story sentences, respectively. ed is the embedding dimension.

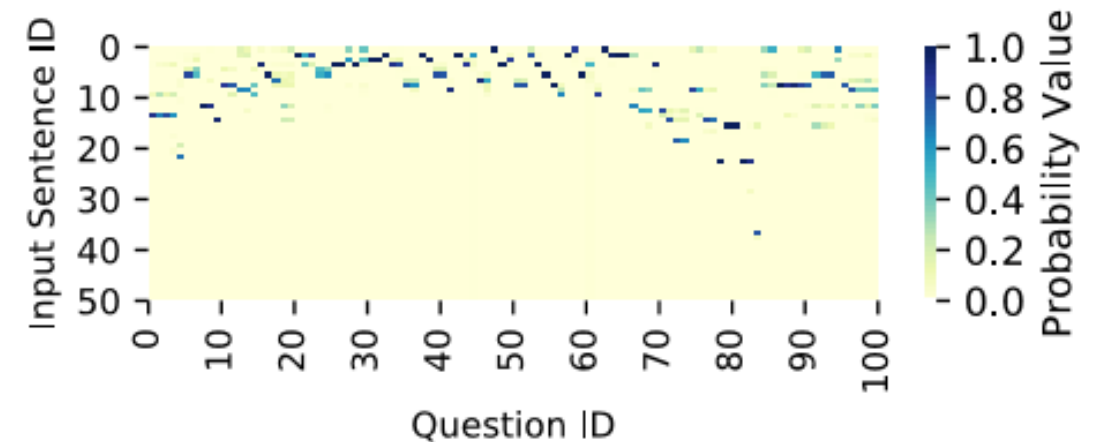


Figure 6: Probability value distribution. Each column represents the probability vector to each question. We use the Facebook bAbi dataset and testset [77].

MnnFast-Zero Skipping

- Algorithm 1 only computes the multiplication when the probability value is larger than a threshold value(th_{skip}).
- There are little trade-off between accuracy and skip ratio.

Algorithm 1: MnnFast's zero-skipping algorithm.

```

input  : The skip threshold  $th_{skip}$ 
input  : The probability vector  $P$ 
input  : The output memory  $M_{IN}$ 
input  : The number of story sentences  $ns$ 
output : The weighted sum  $O$ 

/* Calculate the weighted sum of the output memory with the probability
   values. */
1  $O = [0]$  /* Initialize the output vector. */
2 foreach  $i < ns$  do
   /*  $ns$  is the number of story sentences. */
3   if  $p_i > th_{skip}$  then
4      $O = O + p_i m_i^{OUT}$ 
5 end
6 return  $O$ 

```

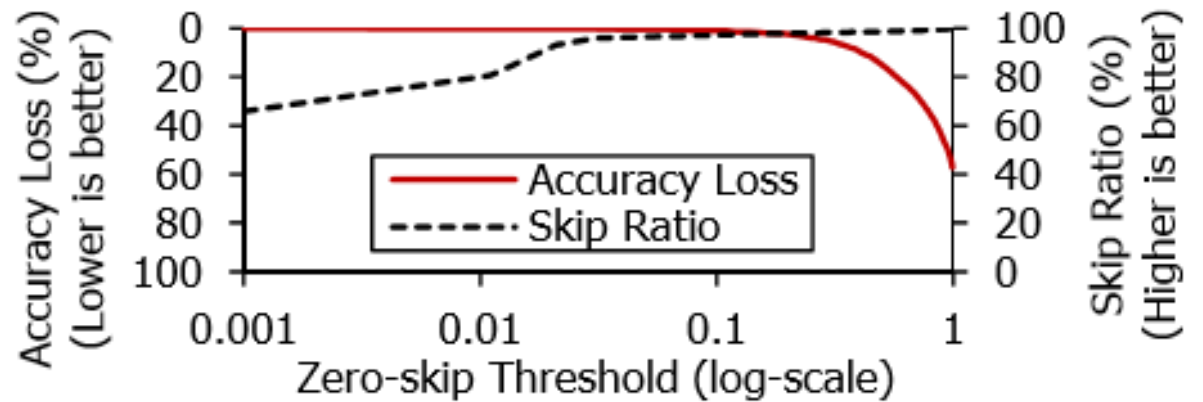


Figure 7: Tradeoffs between accuracy loss (relative loss in accuracy) and computation reduction according to the skip threshold.

MnnFast-Embedding Cache

- **Cache bypassing** has two major drawbacks
 - 1) increases execution latency of embedding operation
 - 2) raises the amount of memory pressure
- **Embedding cache** is a dedicated cache for storing internal state vectors during the embedding operation.

Implementation

- CPU

- ❖ Baseline Implementation

- MemNN is implemented in C++ with open-source BLAS library, OpenBLAS
 - All input/output memory have already been converted into internal data format.
 - Rely on **OpenBLAS** for efficient computation.

- ❖ MNNFast Implementation

- All operation except for softmax are implemented in the **same way to the baseline**.

- GPU

- ❖ GPU kernel implementation

- Rely on **cuBLAS** provided with CUDA toolkit 10.0.
 - Only softmax operation is implemented as one custom kernel.

- ❖ MNNFast Implementation

- Each stream processes chunk and parallelized.
 - Zero-skipping is ineffective

Implementation

- FPGA

- ❖ Baseline Implementation

- Omit the baseline implementation because its design is straightforward

- ❖ MNNFast Implementation

- During embedding, MNNFast converts a question and story sentences by passing them through embedding cache.
 - Embedding cache is designed as a direct mapped cache

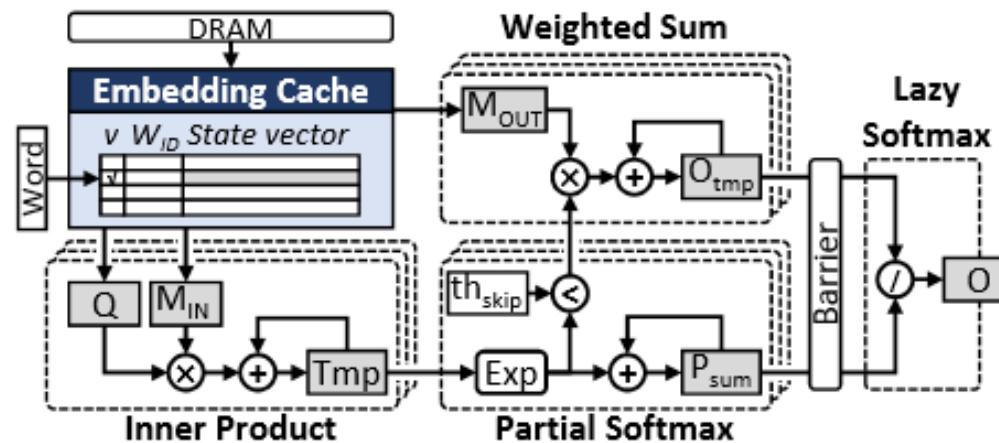




Figure 8: A high-level architecture of FPGA-based MnnFast.

Evaluation

-Performance on CPU



(a) Execution latency breakdown



(b) Performance speedup

Evaluation

- Cache efficiency on CPU




Figure 11: The number of off-chip memory accesses on CPU.





Figure 10: Scalability of column-based algorithm on CPU.

Evaluation

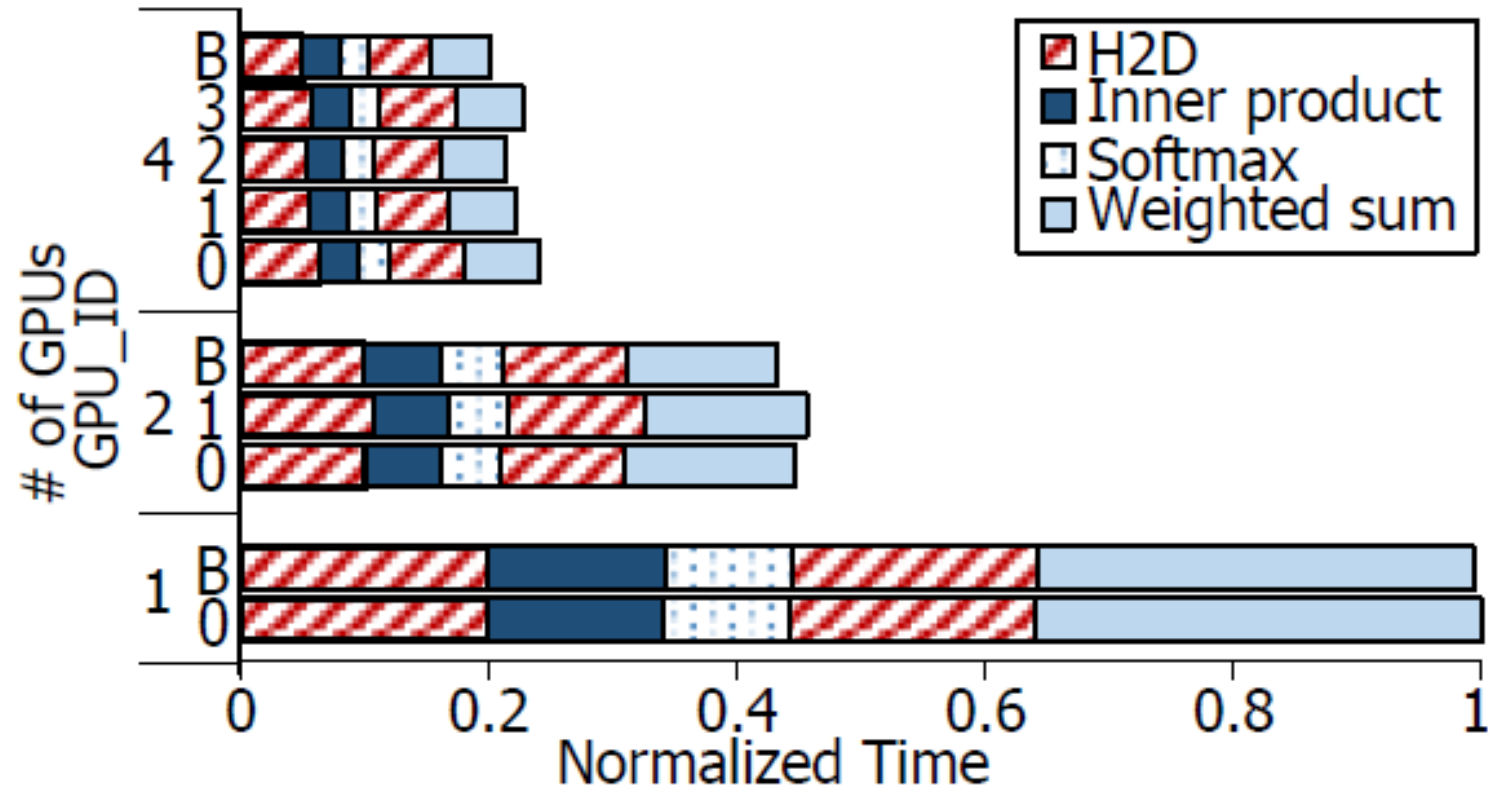
- Multiple CUDA streams



(a) Multiple CUDA streams

Evaluation

- Multiple GPUs



(b) Multiple GPUs (B: best run-alone)

Evaluation

- Performance of FPGA




Figure 13: Latency reduction of FPGA-based MnnFast. Each latency is normalized to the baseline.

Evaluation

- Effectiveness of Embedding Cache in FPGA




Figure 14: Effectiveness of embedding cache in FPGA-based MnnFast. Each latency result is normalized to the No Cache.

Conclusion

- **Three performance problems** of the current architecture :
 - *high memory bandwidth consumption, heavy computation, cache contention.*
- **Three key optimizations** proposed by MnnFast:
 - *column-based algorithm, zero-skipping, and embedding cache.*
- MnnFast solves problem and outperforms the baseline on various hardware : CPU, GPU, and FPGA

Thank you